



# Deterministic Language Architecture: A Formal Category Definition and Architectural Specification for Non-Probabilistic Language Generation Systems

L-SOC — The Lume Synthetic Organism Canon · Language Architecture Volume I

Ronald “Jason” Andrews

*DarkWave Studios LLC*

*Nashville, Tennessee, United States*

*ORCID: 0009-0007-5214-649X*

Version 1.0 — May 2026

**Patent Pending — DarkWave Studios LLC**

**Foundation:** Lume · DOI: [10.5281/zenodo.19382282](https://doi.org/10.5281/zenodo.19382282) · **Trust Layer:** DOI: [10.5281/zenodo.19560674](https://doi.org/10.5281/zenodo.19560674) ·

**Lume-V:** DOI: [10.5281/zenodo.19645097](https://doi.org/10.5281/zenodo.19645097) · **Lume-X:** DOI: [10.5281/zenodo.19443968](https://doi.org/10.5281/zenodo.19443968)

**AXIOM DLA Engine:** [axiom42.com](https://axiom42.com) (reference implementation)

*Term coined:* **DLA — Deterministic Language Architecture** · Coined by Jason Andrews / DarkWave Studios LLC · May 2026

Preprint v1 — Not peer-reviewed. Submitted for early dissemination.

## Abstract

This paper formally defines Deterministic Language Architecture (DLA) — a new category of language generation system in which output is governed by deterministic architecture rather than probabilistic inference. A DLA produces responses by composing structured knowledge through deterministic grammar, synonym selection, and tone governance. It contains no stochastic components in its response generation pipeline. Same input always produces same output. Hallucination is not a risk to be mitigated — it is structurally impossible.

DLA is introduced as a category distinct from Large Language Models (LLMs), expert systems, knowledge-based systems, and neural-symbolic hybrid architectures. The paper provides: a formal architectural definition specifying the necessary and sufficient conditions for a system to qualify as a DLA; a formal proof of the hallucination impossibility property; a comparison of DLA properties against LLMs and prior non-probabilistic language systems; the governance layer argument establishing why a Lume 4/42 synthetic organism is the appropriate architecture for DLA governance; a specification of the auditability and safety properties that follow structurally from DLA architecture; and a description of AXIOM as the reference implementation of DLA.

DLA addresses a gap in the taxonomy of language systems. The field currently distinguishes neural from symbolic, generative from retrieval-based, and fine-tuned from few-shot systems — but has no formal term for the category of systems whose outputs are provably bounded by a verified knowledge base through a deterministic composition function. DLA fills this gap.

**Keywords:** Deterministic Language Architecture, DLA, hallucination impossibility, non-probabilistic language generation, deterministic composition, Lume 4/42, synthetic organism governance, knowledge-bounded generation, formal category definition, AXIOM, auditability, reproducibility, AI safety, natural language generation, compositional grounding

# 1. Introduction

---

## 1.1 The Problem with the Current Taxonomy

The term “language model” has come to mean, in common usage, a large neural network trained on vast text corpora that generates responses by predicting token sequences probabilistically. This conflation of “language system” with “probabilistic neural language system” is technically imprecise and practically consequential.

The imprecision matters because it leaves no shared terminology for language systems that are not probabilistic neural networks. When a practitioner or researcher builds a system that generates natural language responses from a structured knowledge base using deterministic composition logic, there is no accepted category name for what they have built. It is not an LLM. It is not quite an expert system in the traditional sense. It is not a retrieval-augmented generation system. It occupies a well-defined architectural position that the current taxonomy does not name.

The practical consequence of this naming gap is that systems with fundamentally different safety, auditability, and behavioral properties are compared under the same general heading of “AI language systems” — leading to category errors in safety evaluation, regulatory comparison, and user expectation setting.

This paper names the gap. A Deterministic Language Architecture (DLA) is a language system whose outputs are provably bounded by a verified knowledge base through a deterministic composition function, containing no stochastic components in the response generation pipeline.

## 1.2 Why DLA Now

The context for this definition is a global reckoning with the limitations of probabilistic language models. LLMs produce hallucinations — confident, fluent, incorrect statements — at rates that vary by task and model but cannot be reduced to zero by any known training or alignment technique. They produce different outputs for identical inputs depending on sampling temperature and random seed. Their internal reasoning processes are not inspectable. Their training data provenance is often uncertain.

For many applications — legal information, medical guidance, financial advice, scientific reference, safety-critical instructions — these properties are disqualifying. A system that sometimes makes things up is not appropriate for contexts where factual accuracy is

required. A system whose outputs vary run-to-run cannot be audited for compliance. A system whose reasoning is opaque cannot be certified.

DLA provides a formal architecture category for systems that are appropriate for these contexts: systems where hallucination is structurally impossible, outputs are reproducible, reasoning is inspectable, and knowledge provenance is documented.

### 1.3 Contributions

This paper contributes:

- The formal definition of Deterministic Language Architecture as a category — necessary and sufficient architectural conditions
- The hallucination impossibility theorem for DLA systems — a formal proof that hallucination cannot occur in any system satisfying the DLA definition
- A taxonomic comparison of DLA against LLMs, expert systems, knowledge-based systems, retrieval-augmented generation systems, and neural-symbolic hybrid architectures
- The Lume 4/42 organism as the appropriate governance architecture for DLA — the argument for why organism governance produces stronger behavioral guarantees than rule-based governance
- A formal specification of the auditability, safety, and reproducibility properties that follow structurally from DLA architecture
- AXIOM as the reference implementation — the first publicly documented DLA system
- The category name itself — coined May 2026 by Jason Andrews / DarkWave Studios LLC

## 2. Formal Definition

---

### 2.1 The DLA Definition

**Definition 1 (Deterministic Language Architecture):** A language generation system  $L$  is a Deterministic Language Architecture if and only if it satisfies all of the following conditions:

**D1 — Bounded Knowledge:**  $L$  operates from a knowledge base  $K$  that is finite, explicitly enumerated, and human-verifiable.  $K$  contains no implicit knowledge derived from statistical patterns over unenumerated training data.

**D2 — Deterministic Composition:**  $L$ 's response generation function  $f: Q \times K \rightarrow R$  (where  $Q$  is the query space and  $R$  is the response space) is a deterministic pure function. For all  $q \in Q$ ,  $f(q, K)$  returns the same element of  $R$  on every invocation.

**D3 — No Stochastic Components:**  $L$  contains no stochastic components in its response generation pipeline. Pseudo-random elements are permitted only if seeded deterministically from the input (hash-seeded), such that the same input always produces the same pseudo-random sequence.

**D4 — Compositional Grounding:** Every element of every response produced by  $L$  is traceable to a specific element of  $K$ .  $L$  cannot produce content not derivable from  $K$  through its composition function.

**D5 — Inspectable Governance:**  $L$ 's decision logic — how inputs are mapped to outputs — is fully inspectable as a readable function. There are no learned weights, latent embeddings, or attention patterns whose internal representation cannot be read and understood by a human auditor.

A system that satisfies D1–D5 is a DLA. A system that violates any one of D1–D5 is not.

### 2.2 Necessary and Sufficient Conditions

The five conditions are both necessary and sufficient:

**Necessity of D1:** Without a bounded, enumerated knowledge base, the system may draw on implicit statistical knowledge not subject to human verification. A system with unbounded implicit knowledge cannot provide provenance for its outputs.

**Necessity of D2:** Without a deterministic composition function, the system may produce different outputs for identical inputs. Non-deterministic outputs cannot be audited for compliance — the auditor cannot reproduce the specific response that was given.

**Necessity of D3:** Stochastic components in the response pipeline introduce variance that violates D2. Hash-seeded pseudo-randomness is permitted because it preserves the input-output determinism property: the pseudo-random sequence is uniquely determined by the input.

**Necessity of D4:** Without compositional grounding, the system may produce interpolations or extrapolations beyond the content of  $K$ . Compositional grounding is the architectural mechanism that makes hallucination impossible.

**Necessity of D5:** Without inspectable governance, the system’s decision logic cannot be audited, certified, or explained. Opaque governance is a form of non-determinism at the reasoning level even when outputs happen to be reproducible.

**Sufficiency:** A system satisfying D1–D5 has bounded, verified knowledge; deterministic composition; no stochastic elements; provenance for every output element; and inspectable reasoning. No additional conditions are required to produce the hallucination impossibility and reproducibility properties defined in Section 3.

### 3. The Hallucination Impossibility Theorem

---

#### 3.1 Defining Hallucination

**Definition 2 (Hallucination):** A language system  $L$  hallucinates when it produces a response  $r$  containing an assertion  $a$  such that  $a$  is not supported by  $L$ 's knowledge base  $K$  and is presented as factual.

Hallucination has two components: the assertion is not in  $K$ , and it is presented as if it were known. A system that acknowledges uncertainty (“I don’t know”) or explicitly attributes statements to speculation is not hallucinating in the technical sense of this definition.

#### 3.2 The Theorem

**Theorem 1 (Hallucination Impossibility for DLA):** Let  $L$  be a Deterministic Language Architecture satisfying D1–D5. Then  $L$  cannot hallucinate under Definition 2.

**Proof:**

Let  $K$  be the knowledge base of  $L$  (finite, enumerable by D1). Let  $f: Q \times K \rightarrow R$  be  $L$ 's composition function (deterministic pure function by D2, no stochastic components by D3). Let  $r = f(q, K)$  be the response produced by  $L$  for query  $q$ .

By D4 (Compositional Grounding), every element of  $r$  is traceable to a specific element of  $K$ . Formally:

$$\forall \text{ assertion } a \in r, \exists k \in K \text{ such that } a \text{ is derived from } k \text{ by } f$$

Now suppose, for contradiction, that  $L$  hallucinates — that  $r$  contains an assertion  $a$  such that  $a$  is not supported by any  $k \in K$ .

By D4,  $a$  must be derived from some  $k \in K$  by  $f$ . But by assumption, no  $k \in K$  supports  $a$ . This is a contradiction.

Therefore, no response  $r$  produced by  $L$  can contain an assertion not supported by  $K$ .

Furthermore, since  $f$  is deterministic (D2) and  $K$  is verified (D1), the absence of hallucination holds on every invocation, not merely in expectation. ■

### 3.3 The Contrast with Probabilistic Systems

In a probabilistic language model  $M$ , the response is generated by sampling from a conditional probability distribution:

$$r \sim P(R \mid q, \theta)$$

Where  $\theta$  represents the model’s learned parameters. The distribution  $P(R \mid q, \theta)$  is defined over the entire response space  $R$ , which is not bounded by any explicit knowledge base.  $P$  assigns non-zero probability to any response that the model’s architecture can represent — including responses containing assertions not present in any training document, not present in any retrieval source, and not grounded in any verified fact.

The expected hallucination rate of  $M$  can be reduced by better training, retrieval augmentation, and constitutional alignment methods, but it cannot be reduced to zero. For any positive temperature  $\tau > 0$ , there exists a non-zero probability of sampling a hallucinatory response. At  $\tau = 0$  (greedy decoding), the system becomes more deterministic but does not become compositionally grounded — the output is still a function of learned weights over unverified training data.

The hallucination impossibility in a DLA is not a probabilistic claim about expected behavior. It is a structural guarantee that follows from the architecture. No configuration of a DLA can hallucinate; no configuration of a probabilistic language model can structurally prevent hallucination.



## 4. Taxonomic Comparison

### 4.1 DLA vs. Large Language Models

| Property           | LLM                                                | DLA                                                          |
|--------------------|----------------------------------------------------|--------------------------------------------------------------|
| Output determinism | Probabilistic (varies by temperature and seed)     | Deterministic (same input $\rightarrow$ same output, always) |
| Knowledge base     | Implicit (learned weights over training corpus)    | Explicit (enumerated, human-verifiable $K$ )                 |
| Hallucination      | Structurally possible; mitigated by training       | Structurally impossible by Theorem 1                         |
| Output provenance  | Not traceable to specific training examples        | Traceable to specific $k \in K$                              |
| Auditability       | Opaque (attention weights not human-interpretable) | Fully inspectable (pure function)                            |
| Infrastructure     | Cloud, GPU, token billing                          | Local, offline-capable                                       |
| Safety enforcement | Guardrails bolted on post-training                 | Hard constraints at architecture level                       |
| Data licensing     | Training data provenance often uncertain           | Explicit, documented, verifiable                             |
| Domain adaptation  | Fine-tuning or RAG                                 | Knowledge base extension                                     |
| Response variance  | High (intentionally varied by temperature)         | Zero (D3: same seed $\rightarrow$ same output)               |

### 4.2 DLA vs. Expert Systems

Expert systems, developed primarily in the 1980s, represent a prior generation of non-probabilistic language systems. They are relevant predecessors to DLA but differ in several significant ways:

| Property                 | Expert System                            | DLA                                           |
|--------------------------|------------------------------------------|-----------------------------------------------|
| Knowledge representation | Production rules (IF-THEN)               | Structured facts + grammar templates          |
| Response generation      | Rule inference chain → formatted output  | Deterministic NLG composition                 |
| Natural language quality | Formulaic, rule-formatted                | Natural, varied (deterministically varied)    |
| Governance architecture  | Forward/backward chaining inference      | Organism-governed multi-primitive system      |
| Scale                    | Typically hundreds to thousands of rules | Hundreds of thousands of topics               |
| Synonym variation        | None — identical phrasing every response | Deterministic synonym selection (hash-seeded) |

A DLA is not an expert system. Expert systems use explicit production rules to derive conclusions; a DLA uses structured knowledge facts and deterministic grammar to compose natural language responses. Expert systems typically do not satisfy D4 in the compositional grounding sense — expert systems can derive conclusions that are logical implications of their rules rather than direct instantiations of enumerated facts, and those implications may be incorrect if the rule base is incomplete or inconsistent.

### 4.3 DLA vs. Knowledge-Based Systems

Knowledge-based systems (KBS) use structured knowledge representations (ontologies, semantic networks, description logics) to answer queries. They are closer to DLA than expert systems but differ in a key way: most KBS are query-answer systems, not natural language generation systems. They retrieve or derive answers but do not compose them into natural language text through a deterministic NLG pipeline. A DLA extends KBS with a deterministic NLG layer that satisfies D2 and D4.

### 4.4 DLA vs. Retrieval-Augmented Generation (RAG)

RAG systems combine a retrieval engine (which fetches relevant documents from a corpus) with a neural language model (which generates responses conditioned on the retrieved documents and the query). RAG systems reduce hallucination by grounding generation in retrieved content, but they do not eliminate it:

- The neural generation component retains its probabilistic nature — it can generate content not present in the retrieved documents
- Retrieved documents may themselves contain incorrect information
- The system is not deterministic — identical queries may retrieve different documents and the generation component introduces sampling variance

A RAG system violates D2 (not deterministic), D3 (stochastic generation component), and D4 (no compositional grounding guarantee). It is not a DLA.

#### 4.5 DLA vs. Neural-Symbolic Hybrid Systems

Neural-symbolic systems combine neural components (for perception, embedding, or language understanding) with symbolic components (for reasoning, planning, or knowledge representation). The critical limitation from the DLA perspective is that the presence of any neural component in the response generation pipeline violates D3 and typically D2. Even if the symbolic component is fully deterministic, the neural component introduces the probabilistic properties — and potential hallucination — of purely neural systems. A neural-symbolic system reduces hallucination relative to a pure LLM but cannot satisfy the DLA definition.

## 5. The Governance Layer Argument

---

### 5.1 Why Rule-Based Governance Is Insufficient

Previous non-probabilistic language systems — expert systems and KBS — use rule-based governance: explicit IF-THEN conditions that determine system behavior. Rule-based governance is deterministic and inspectable (satisfying D2 and D5) but has a structural limitation: rules interact.

When two rules fire simultaneously and their conclusions conflict, a conflict resolution mechanism is required. When a rule applies in a context its author did not anticipate, the rule may produce an incorrect conclusion. When the rule base grows large, interactions between rules become difficult to predict and verify. Rule-based systems degrade gracefully only when their rule bases are carefully curated and conflict-free — a maintenance challenge that grows superlinearly with the number of rules.

### 5.2 Organism Governance as the DLA Standard

A Lume 4/42 synthetic organism provides a governance architecture that addresses the rule interaction problem through a fundamentally different approach: instead of encoding behavior as explicit rules that fire conditionally, the organism encodes behavior as a continuous normalized state space with discrete mode selection.

The organism's governance logic is:

1. Normalize all inputs to  $[-1.0, +1.0]$
2. Compute primitive aggregates
3. Classify each primitive against thresholds (optimal / advisory / caution / critical)
4. Select exactly one operating mode via a priority hierarchy
5. Issue coherent, simultaneous outputs to all actuators (or, in the language context, all composition parameters)

This architecture has no rule interactions — there is no conflict resolution problem because there are no rules to conflict. There is only one mode at any moment, selected by a transparent priority hierarchy. The same inputs always select the same mode; the same mode always produces the same composition parameters; the same composition parameters always produce the same response.

### 5.3 The Organism as DLA Governor

In a DLA governed by a Lume 4/42 organism, the organism controls the composition parameters that determine how the knowledge fact is rendered into language:

```

Query → Intent Classification → Knowledge Engine (K lookup)
  → Structured Fact { subject, core, process, goal, aspects }
  → Organism Mode Selection:
    PRECISION MODE: technical tone, minimal synonym variation,
maximum fact fidelity
    CLARITY MODE: conversational tone, moderate variation,
accessible vocabulary
    EMPATHY MODE: empathetic tone, softer variation, emotionally
appropriate framing
    CONCISE MODE: brief templates, minimal elaboration
    SCHOLARLY MODE: formal register, citation-appropriate framing
  → Grammar template selected by mode + intent
  → Synonym variation applied (hash-seeded from query)
  → Response composed

```

The organism is not determining what facts to present — that is determined by the knowledge engine lookup. The organism is determining *how* to present them — the register, the template complexity, the synonym selection depth. This is a governance function, not a knowledge function. The organism governs presentation; the knowledge base governs content.

This separation of concerns is architecturally clean and satisfies D4: even with organism-governed presentation variation, every element of the response remains traceable to a specific  $k \in K$ . The organism can change how a fact is expressed; it cannot introduce content not in  $K$ .

### 5.4 Hard Constraints in DLA

The Lume organism architecture includes hard constraints: conditions that trigger mandatory responses regardless of mode selection, that cannot be overridden by any higher-level logic. In the DLA context, hard constraints govern the integrity of the response pipeline:

**Hard constraint 1 — Knowledge boundary:** If a query cannot be matched to any  $k \in K$ , the system must respond with an explicit acknowledgment of unknown status. It cannot compose a response from outside  $K$ . This is D4 enforced as a hard constraint.

**Hard constraint 2 — Correction integrity:** If a user correction conflicts with a  $k \in K$  that carries a high confidence marker, the correction is logged but  $K$  is not immediately overwritten. The conflict is surfaced for human review. This protects D1 against casual corruption.

**Hard constraint 3 — Response traceability:** Every response includes a traceable provenance path from output to source  $k \in K$  available on request. The system can never produce a response whose provenance it cannot provide.

## 6. Auditability, Safety, and Reproducibility

---

### 6.1 Auditability

A DLA is fully auditable by construction. Given a query  $q$  and a response  $r$ , an auditor can:

1. Identify the intent classification applied to  $q$
2. Identify the specific  $k \in K$  that was retrieved for  $q$
3. Identify the grammar template applied to  $k$
4. Identify the synonym substitutions applied (from the hash-seeded synonym selection)
5. Identify the organism mode that governed the presentation
6. Verify that every element of  $r$  is derivable from  $k$ , the template, and the synonym table

This audit trail is a pure function trace — the same reasoning that produced  $r$  will produce  $r$  again. No LLM provides this auditability. The reasoning behind an LLM response is distributed across billions of learned weights and cannot be reconstructed from the output alone.

### 6.2 Regulatory Compliance

The auditability property makes DLA systems appropriate for regulatory compliance in ways that LLMs structurally are not:

**Financial services:** Regulators require that advice given to clients be reproducible and traceable to approved information sources. A DLA’s provenance chain satisfies this; an LLM’s cannot.

**Medical information:** Health information regulations in multiple jurisdictions require that AI-provided health information be traceable to approved clinical sources. A DLA’s D1 satisfies this if  $K$  consists of approved clinical content; an LLM’s implicit knowledge does not.

**Legal information:** Bar rules on unauthorized practice of law increasingly address AI legal information systems. A DLA that is clear about the boundaries of its knowledge satisfies the transparency requirements that are emerging; an LLM that may confabulate legal holdings does not.

### 6.3 Safety Properties

**Safety Property 1 — Bounded harm from knowledge errors:** If a  $k \in K$  is incorrect, the harm from that error is bounded to queries that match  $k$ . In an LLM, a problematic training example can influence responses across many topics through learned associations. In a DLA, an incorrect  $k$  affects only queries that retrieve  $k$ . Harm is localized and correctable.

**Safety Property 2 — Predictable failure modes:** When a DLA fails, the failure mode is predictable: either  $K$  does not contain relevant content or  $k \in K$  is incorrect. These failure modes are auditable and correctable.

**Safety Property 3 — No emergent harmful content:** A DLA cannot produce harmful content not present in  $K$ . If  $K$  is curated to exclude harmful content, no harmful content can appear in responses.

### 6.4 Reproducibility

A DLA satisfies the scientific reproducibility requirement: the same inputs produce the same outputs, always, without exception. A response given to user A today will be given to user B tomorrow under identical query conditions. A response given in a regulatory audit can be reproduced exactly during litigation. These reproducibility properties are a direct consequence of D2 and D3.



## 7. AXIOM — The Reference Implementation

### 7.1 System Overview

AXIOM ([axiom42.com](http://axiom42.com)) is the first publicly documented Deterministic Language Architecture. It satisfies all five DLA conditions:

**D1 — Bounded Knowledge:** AXIOM operates from a knowledge base of approximately 181,000 topics across 212 domain packs. Every topic is explicitly authored, human-verified, and stored as structured text.

**D2 — Deterministic Composition:** AXIOM's response generation pipeline is a deterministic pure function: query  $\rightarrow$  intent classification  $\rightarrow$  knowledge engine lookup  $\rightarrow$  fact decomposition  $\rightarrow$  grammar template selection  $\rightarrow$  synonym application  $\rightarrow$  response.

**D3 — No Stochastic Components:** AXIOM contains no neural networks, probability distributions, or sampling operations in its response pipeline. Synonym variation uses hash-seeded pseudo-random selection.

**D4 — Compositional Grounding:** Every element of every AXIOM response is derived from a specific topic in the knowledge base through the composition function.

**D5 — Inspectable Governance:** AXIOM's composition engine, grammar engine, tone adapter, and thesaurus are all inspectable as source code.

### 7.2 The NLG Pipeline as DLA Implementation

```
Input: query q, knowledge base K
  → Intent classification: classify(q) → intent ∈ {DEFINE, EXPLAIN,
SUMMARIZE, ...}
  → Knowledge lookup: lookup(q, K) → k = { subject, core, process,
goal, aspects, keywords }
  → Grammar selection: select_template(intent) → template ∈ G
  → Fact instantiation: instantiate(template, k) → draft_response
  → Synonym variation: vary(draft_response, hash(q), thesaurus) →
varied_response
  → Tone application: apply_tone(varied_response, detect_tone(q)) →
final_response
Output: final_response r ∈ R
```

This pipeline is a deterministic pure function. No step introduces stochasticity.

### 7.3 Organism Governance in AXIOM

AXIOM’s tone adapter implements a simplified organism governance layer — a direct application of the organism principle: discrete mode selection from a normalized state evaluation. The tone detection function evaluates the query against a priority hierarchy:

1. Learned user preference (persistent from prior sessions)
2. Explicit query signal (keywords indicating register preference)
3. Domain default (domain-specific tone mapping)
4. Conversational fallback

The full Lume 4/42 organism architecture is the appropriate governance layer for a DLA at scale — governing not just tone but composition depth, synonym variation rate, elaboration length, and cross-query context coherence simultaneously.

### 7.4 Learning System Within DLA Constraints

AXIOM’s five-dimension learning system (alias learning, fact learning, correction learning, preference learning, conversation context) operates within the DLA constraints. New facts enter  $K$  as explicitly authored entries, satisfying D1. Corrections modify  $k$  entries subject to the correction confidence layer. Preferences govern organism mode selection. The learning system extends  $K$  while preserving D1–D5.

## 8. Positioning DLA in the AI Safety Landscape

### 8.1 The Alignment Problem and DLA

The AI alignment problem — ensuring that AI systems behave in accordance with human intentions and values — is one of the most active research areas in artificial intelligence. DLA offers a different relationship to the alignment problem: not solving alignment in a neural system, but building a system in which certain categories of misalignment are architecturally impossible.

A DLA cannot produce content outside its knowledge base (D4). It cannot behave non-deterministically across users or sessions (D2, D3). It cannot have hidden reasoning that differs from its inspectable governance logic (D5). These are not alignment properties achieved through training — they are structural properties of the architecture.

DLA does not claim to solve the alignment problem in general. It claims to be architecturally immune to specific failure modes — hallucination, output variance, opaque reasoning — that are alignment concerns in neural systems.

### 8.2 The Safety Case for DLA Deployment

When evaluating whether a language system is appropriate for deployment in a safety-sensitive context, the relevant questions are:

| Safety Question                                                  | DLA Answer              | LLM Answer                  |
|------------------------------------------------------------------|-------------------------|-----------------------------|
| Can the system produce incorrect factual assertions?             | No (Theorem 1)          | Yes (hallucination)         |
| Can the system produce different responses to identical queries? | No (D2, D3)             | Yes (temperature, sampling) |
| Can the system's decision logic be audited?                      | Yes (D5)                | No (opaque weights)         |
| Can harmful content be excluded architecturally?                 | Yes (D4, curated $K$ )  | Conditionally (filtering)   |
| Can the system's knowledge be updated and verified?              | Yes (D1, explicit $K$ ) | Requires retraining         |

## 9. DLA Qualification Criteria

---

### 9.1 How to Determine if a System Is a DLA

The following checklist operationalizes the formal DLA definition for practical system evaluation:

- ☐ **D1 Check:** Can you enumerate every fact the system may present in a response? Is that enumeration finite and human-readable? If yes: D1 satisfied.
- ☐ **D2 Check:** Given the same query, does the system produce the same response on repeated invocations with no external state changes? Test with 100 identical queries. If all 100 responses are identical: D2 satisfied.
- ☐ **D3 Check:** Does the system's response generation pipeline contain any neural networks, probability distributions, sampling operations, or random number generators not seeded deterministically from the input? If no: D3 satisfied.
- ☐ **D4 Check:** Can you trace every sentence in a generated response to a specific source entry in the knowledge base? If yes for all tested responses: D4 satisfied.
- ☐ **D5 Check:** Can a human auditor read the complete decision logic that determined the response, without reference to learned weights, embeddings, or attention patterns? If yes: D5 satisfied.

A system satisfying all five checks is a DLA. A system failing any check is not.

### 9.2 Partial DLA Systems

Some systems may satisfy some but not all DLA conditions. Partial satisfaction provides partial benefits. A system with D1, D2, D3, D4 but not D5 is close to a DLA and safer than an LLM, but not certifiable as DLA for compliance purposes. The value of the formal definition is precisely that it sets a clear bar.

## 10. Intellectual Property Context

---

### 10.1 The Coined Term

The term “Deterministic Language Architecture” and its abbreviation “DLA” are coined in this paper by Jason Andrews / DarkWave Studios LLC, May 2026. Patent Pending.

Prior art search conducted by the author found no prior use of “Deterministic Language Architecture” as a formal category name in academic literature, industry publications, or patent filings as of the date of this paper. This paper is the prior art anchor for the DLA category.

### 10.2 Novel Claims

- The formal DLA category definition (D1–D5) — the necessary and sufficient conditions for a language system to qualify as a DLA
- The Hallucination Impossibility Theorem for DLA systems
- The Lume 4/42 organism as the governance architecture for DLA
- Hard constraints for DLA integrity
- The DLA qualification checklist

## 11. Discussion and Future Work

---

### 11.1 DLA at Scale

AXIOM demonstrates DLA at 181,000 topics. The formal definition places no upper bound on  $K$  — a DLA can, in principle, have a knowledge base of any size, provided D1 is satisfied. Practical upper bounds are imposed by knowledge authoring capacity rather than by the architecture.

### 11.2 DLA and Organism Governance at Full Scale

The AXIOM reference implementation uses a simplified organism governance layer. The full Lume 4/42 organism architecture — forty-two nodes, four primitives, five operating modes, hard constraints — is the appropriate governance target for a DLA that must simultaneously optimize composition depth, synonym variation, elaboration strategy, contextual coherence, and safety constraint enforcement across a very large knowledge base.

Future work includes specifying the complete 42-node Lume organism mapping for the DLA language domain — an organism that governs language composition the way HydroCore governs hydraulic pressure and HydroCore Steam governs industrial steam.

### 11.3 DLA Certification

A formal certification framework — analogous to ISO standards for software quality or ASME standards for pressure vessels — could establish DLA certification as a recognized compliance status for language systems in regulated industries.

### 11.4 The Ecosystem of Determinism

DLA is the language layer of a broader ecosystem of deterministic governance architectures. The Lume synthetic organism governs physical systems (HydroCore, Meridian), biological systems (BioCore), cognitive systems (NeuroCore), social systems (SocioCore), governance systems (GovernanceCore), and now language generation (DLA). The organism architecture is domain-invariant. The governance logic is universal.

The vision is an ecosystem in which every layer of a system — from physical infrastructure to language interface — is governed by a deterministic organism that provides the same formal guarantees: same inputs, same outputs; inspectable reasoning;

hard constraints enforced at the architecture level; no guessing at any layer.

DLA is the language manifestation of that vision. It closes the circle from physical to language — the same deterministic governance principle that governs a steam turbine governs the words used to describe it.

## 12. Conclusion

---

Deterministic Language Architecture is a necessary category. The gap it fills — a formal name for language systems that are not probabilistic, whose outputs are provably bounded by verified knowledge, and whose reasoning is fully inspectable — has existed since the first LLM was deployed in a context where its hallucination risk was a disqualifying property.

The formal definition (D1–D5) is precise, operationalizable, and admits a formal proof of the hallucination impossibility theorem. The governance layer argument establishes why Lume 4/42 organism governance is the appropriate architecture for DLA systems at scale. The AXIOM reference implementation demonstrates that DLA systems are not theoretical — they exist, they are deployed, and they perform.

The category name is coined. The formal definition is established. The prior art is anchored.

**DLA systems do not guess. They know what they know, they say it the same way every time, and they cannot say what they do not know.**



## Appendix A — DLA Formal Notation Summary

| Symbol                        | Meaning                                                |
|-------------------------------|--------------------------------------------------------|
| $K$                           | Knowledge base (finite, enumerable set of facts)       |
| $Q$                           | Query space (set of all possible queries)              |
| $R$                           | Response space (set of all possible responses)         |
| $G$                           | Grammar template set (finite)                          |
| $S$                           | Synonym table (finite, hash-seeded selection)          |
| $f: Q \times K \rightarrow R$ | DLA composition function (deterministic pure function) |
| $k \in K$                     | A specific knowledge fact                              |
| $q \in Q$                     | A specific query                                       |
| $r \in R$                     | A specific response                                    |
| $a \in r$                     | An assertion contained in response $r$                 |
| $\text{hash}(q)$              | Deterministic hash of query $q$ (synonym seed)         |

## Appendix B — DLA Qualification Checklist (Reproducible)

### DLA QUALIFICATION CHECKLIST

System: \_\_\_\_\_

Date: \_\_\_\_\_

Evaluator: \_\_\_\_\_

#### D1 -- Bounded Knowledge

- ☐ Knowledge base is finite and enumerable
- ☐ Every entry is human-authored or human-verified
- ☐ No implicit statistical knowledge from unenumerated training data

Result: PASS / FAIL

#### D2 -- Deterministic Composition

- ☐ 100 identical queries produce 100 identical responses
- ☐ No external state changes between test runs
- ☐ Result verified by independent evaluator

Result: PASS / FAIL

#### D3 -- No Stochastic Components

- ☐ No neural networks in response generation pipeline
- ☐ No probability distributions in response generation pipeline
- ☐ No random number generators except hash-seeded pseudo-random

Result: PASS / FAIL

#### D4 -- Compositional Grounding

- ☐ 20 sampled responses each traced to specific K entries
- ☐ No response element found without traceable K provenance
- ☐ Hard constraint 1 verified: unknown queries produce explicit unknown response

Result: PASS / FAIL

#### D5 -- Inspectable Governance

- ☐ Complete decision logic readable as source code or specification
- ☐ No learned weights, embeddings, or attention patterns in decision path
- ☐ Independent auditor can reproduce full response derivation

Result: PASS / FAIL

OVERALL DLA QUALIFICATION: PASS (all 5) / FAIL (any 1)

## Appendix C — The Impossibility Proof (Compact Form)

---

Let  $L$  be a DLA (satisfies D1-D5).  
Let  $K$  be  $L$ 's knowledge base (finite, D1).  
Let  $f: Q \times K \rightarrow R$  be  $L$ 's composition function (deterministic, D2; no stochastic, D3).  
Let  $r = f(q, K)$  for any query  $q$ .  
By D4: for all  $a$  in  $r$ , there exists  $k$  in  $K$ :  $a$  is derived from  $k$  by  $f$ .  
Suppose  $L$  hallucinates: there exists  $a$  in  $r$  such that no  $k$  in  $K$  supports  $a$ .  
By D4: there exists  $k$  in  $K$ :  $a$  is derived from  $k$ . Contradiction.  
Therefore:  $L$  cannot hallucinate. QED

## References

---

- [1] J. Andrews, “LUME: Eliminating Cognitive Distance,” DarkWave Studios LLC, 2026. DOI: [10.5281/zenodo.19382282](https://doi.org/10.5281/zenodo.19382282). Patent Pending.
- [2] J. Andrews, “Lume-V: The Deterministic Wrapper,” DarkWave Studios LLC, 2026. DOI: [10.5281/zenodo.19645097](https://doi.org/10.5281/zenodo.19645097). Patent Pending.
- [3] J. Andrews, “Lume-X: The Multi-Organism Substrate,” DarkWave Studios LLC, 2026. DOI: [10.5281/zenodo.19443968](https://doi.org/10.5281/zenodo.19443968). Patent Pending.
- [4] J. Andrews, “HydroCore Physical,” DarkWave Studios LLC, 2026. L-SOC Series. Patent Pending.
- [5] J. Andrews, “HydroCore Drive,” DarkWave Studios LLC, 2026. L-SOC Series. Patent Pending.
- [6] J. Andrews, “HydroCore Steam,” DarkWave Studios LLC, 2026. L-SOC Series. Patent Pending.
- [7] J. Andrews, “Meridian Infrastructure,” DarkWave Studios LLC, 2026. L-SOC Series. Patent Pending.
- [8] E. M. Bender, T. Gebru, A. McMillan-Major, and S. Shmitchell, “On the dangers of stochastic parrots: Can language models be too big?” *Proceedings of FAccT 2021*, pp. 610–623.
- [9] G. Marcus, “Deep learning is hitting a wall,” *Nautilus*, MIT Press, 2022.
- [10] J. Maynez, S. Narayan, B. Bohnet, and R. McDonald, “On faithfulness and factuality in abstractive summarization,” *Proceedings of ACL 2020*, pp. 1906–1919.
- [11] L. Huang et al., “A survey on hallucination in large language models,” *arXiv preprint arXiv:2311.05232*, 2023.
- [12] Z. Ji et al., “Survey of hallucination in natural language generation,” *ACM Computing Surveys*, vol. 55, no. 12, pp. 1–38, 2023.
- [13] D. B. Lenat and R. V. Guha, *Building Large Knowledge-Based Systems*, Addison-Wesley, 1990.
- [14] P. Jackson, *Introduction to Expert Systems*, 3rd ed., Addison-Wesley, 1998.
- [15] E. Reiter and R. Dale, *Building Natural Language Generation Systems*, Cambridge University Press, 2000.

---

**Patent Pending** — DarkWave Studios LLC.

Term “Deterministic Language Architecture” and abbreviation “DLA” coined by Jason Andrews / DarkWave Studios LLC, May 2026.

© 2026 DarkWave Studios LLC. All rights reserved.

Ronald “Jason” Andrews · DarkWave Studios LLC · Nashville, TN · [team@dwsc.io](mailto:team@dwsc.io)

ORCID: [0009-0007-5214-649X](https://orcid.org/0009-0007-5214-649X) · [lume-lang.org](https://lume-lang.org) · [github.com/cryptocreper94-sudo](https://github.com/cryptocreper94-sudo)

This preprint has not undergone peer review. Submitted for early dissemination and to establish priority of invention.